



UNIVERSITY OF
LIVERPOOL

DEPARTMENT OF COMPUTER SCIENCE · COMP702 MSC PROJECT

Kodro: An Offline, Self-Refining Platform for Designing Robots and Validating Their Behaviour in a Realistic Simulation

*Dissertation submitted in partial fulfilment of the requirements for the degree of
Master of Science*

Author Vaibhav Lalwani

Student ID 201979723

Supervisor Keith Dures

Submission September 2026

Kodro: An Offline, Self-Refining Platform for Designing Robots and Validating Their Behaviour in a Realistic Simulation

Anonymous title page for marking

Student ID 201979723

Acknowledgements

I thank my supervisor, Keith Dures, for his guidance throughout this project.

To complete: personal acknowledgements before final submission.

Abstract

This dissertation presents Kodro, an offline desktop platform on which a non-expert designs a custom robot, programs it by code, blocks or voice, and validates how it behaves in a realistic simulated world before committing to real hardware. The work studies a focused question: how far an artificial intelligence assistant can usefully self refine from accumulated use, and how trustworthy a simulator can be made, under a hard constraint that everything runs locally with no cloud service, no accounts and no retraining of the model. The platform grounds every suggestion in the robot's own capability specification so the assistant cannot propose a part the build does not have, reviews the generated control code, and validates it across randomised conditions in worlds that range from a street with pedestrians and traffic to an indoor room. Self refinement is delivered honestly at the level of the system, through reflection memory and a growing skill library rather than weight retraining, which remains unsolved offline. The dissertation reports what was built, how it was evaluated against usability, design and self refinement, and what the results say about the offline boundary.

To finalise after results: add the measured headline numbers (design study task success, self refinement trend, AI correctness and invention rates) once the evaluation has run.

Statement of Ethical Compliance

The declared categories for this project are data category A0 and participant category 2. I confirm that I have followed the School's ethical guidance throughout. The development and the simulated persona evaluation involved no human participants and no personal data. The planned design study with human participants was conducted only with appropriate consent and within the activities permitted under the agreed categories, and any data collected was anonymised at the point of collection and is not linked to any individual.

Table of Contents

- 1 Introduction and Background
- 2 Requirements Analysis
- 3 Design
- 4 Implementation

5 Testing and Evaluation

6 Project Ethics

7 Conclusion and Future Work

References

Appendices

Note: page numbers are added on the final typeset pass. Maximum 50 pages excluding the appendix, no word count, PDF only, checked against Turnitin. Due 11 September 2026.

1. Introduction and Background

1.1 Motivation

Plenty of people have an idea for a robot and no safe way to try it. Buying parts, wiring a board and writing control code is slow and expensive, and the first time a real robot meets the real world it tends to drive into something. Kodro closes that gap with a single offline download. A user designs a robot in a builder, choosing a chassis, a controller board such as an ESP32 or a micro:bit, and the sensors and motors it carries; programs its behaviour in real Python, in blocks, or by voice, with grounded artificial intelligence help and a second agent that reviews the code; then runs it in a simulated world that behaves like the real one, with terrain, physics and other agents such as pedestrians and traffic moving around it.

1.2 Aim and research question

The project asks how much an offline system can genuinely refine itself, and how trustworthy its validation can be, under a hard constraint that rules out the cloud and rules out changing the model's weights. This dissertation describes the platform that was built to study that question and reports what the study found.

1.3 Background

The design rests on three bodies of work. Constructionism holds that people learn most when they build something they can watch behave and then debug (Papert, 1980), while cognitive load theory warns that an elaborate scene can spend a beginner's attention on the view rather than the problem (Sweller, 1988). The central problem in robotics is the sim to real gap, the drop in performance when behaviour tuned in simulation meets the real world, for which the standard answer is domain randomization rather than a perfect simulator (Tobin et al., 2017); realistic worlds also need agents that move with intent, drawn here from the social force model (Helbing and Molnar, 1995) and reciprocal velocity obstacles (van den Berg et al., 2008). The assistant builds on grounding generation in retrieved sources to control hallucination (Lewis et al., 2020; Huang et al., 2023) and on programs written against a fixed robot interface (Liang et al., 2023; Ahn et al., 2022), while its self refinement follows verbal reflection (Shinn et al., 2023) and a growing skill library (Wang et al., 2023) rather than the unsolved weight level evolution (Tao et al., 2024).

1.4 Where Kodro sits among existing tools

Kodro is not the first tool to let people build and simulate robots, and it is worth being clear about what is and is not new. Professional robot simulators such as Gazebo, Webots and Isaac Sim are powerful but heavy, assume a robotics background, and lean on cloud or workstation resources; they are aimed at researchers, not a beginner on a school laptop. Visual and block tools such as Scratch, Blockly and the VEX and Lego ecosystems lower the entry cost and are rightly common in classrooms, but they stop short of validating a custom design in a realistic, agent populated world. End user robotics has long wrestled with the same tension between approachability and capability. Kodro sits deliberately in the gap: it keeps the approachable design and block path of the classroom tools, adds the realistic validation of the heavy simulators

in a form light enough to run offline on one machine, and binds a grounded assistant to the user's own build. The novelty is not any single technique but the combination, and the honest study of how far it can be pushed under a strict offline constraint.

1.5 Contributions and structure

The contributions of this work are the design and implementation of that offline platform; an honest, system level account of self refinement that is implemented and observed rather than claimed; a unified moving agent simulation that the robot must genuinely cope with; and an evaluation design that separates usability, learning and self refinement into measurable questions. The remainder of the dissertation is structured as follows. Chapter 2 sets out the requirements; Chapter 3 the design; Chapter 4 the realised implementation; Chapter 5 the testing and evaluation, including the results in hand and the studies still to run; Chapter 6 the project ethics; and Chapter 7 the conclusions and future work.

2. Requirements Analysis

The requirements are stated with a plain test of done so progress can be judged rather than asserted. The essential requirements are: a robot design whose specification drives the simulation; sandboxed execution of the robot's real Python validated against per scenario success criteria; validation in a realistic world across randomised conditions; an artificial intelligence layer that runs on a local model only with a deterministic rule based fallback; and every byte of user data kept on the device.

Requirement	Test of done
Specification drives the simulation	Changing the build measurably changes behaviour: a heavier robot drains its battery faster, a stronger motor set raises top speed, a fitted sensor unlocks its command and an absent one withholds it.
Sandboxed, validated execution	A hostile program cannot reach the disk and a correct program is scored correctly against the scenario.
Realistic, repeatable validation	The same program runs across several randomised seeds and a per run report of successes, collisions and time to goal is produced.
Local only AI with fallback	The platform is fully usable with the network switched off and no model present, falling back within two seconds.
Data stays on the device	A network trace during normal use shows no traffic.

2.1 Functional requirements

The functional requirements describe what the platform does, each traced to the part of the implementation that meets it.

The platform shall	Met by
Let a user design a custom robot from a chassis, a controller board, sensors and actuators	The Robot Lab, where the chosen parts produce a single specification object the rest of the system reads.
Let the user program the robot by text, blocks or voice	A Python editor, a block palette that compiles to the same Python, and a voice route that turns speech into the same commands.
Assist with code without inventing capabilities the build lacks	The grounded assistant and a second agent reviewer, both reading the robot specification.
Run the robot in a realistic world and report behaviour	The simulation in three dimensional worlds with physics, terrain and moving agents, returning successes, collisions and time to goal.
Refine its help from accumulated use	The local memory store: per run reflections, a reusable skill library, and an experience table.

2.2 Non functional requirements

Quality	Requirement and test
Offline and private	No outbound network request in normal use; a network trace shows no traffic and all data stays in one local file.
Low resource	Runs on a mainstream laptop with no discrete graphics card and an eight gigabyte memory floor.
Robust	A failed graphics context or a missing model degrades to a working fallback, never a blank screen or a crash.
Accessible	WCAG 2.1 AA: keyboard only operation, the AA contrast ratio, reduced motion honoured, read aloud and high contrast options.
Maintainable	A continuous integration gate runs the full test suite, a type checker and a linter on three operating systems before any change merges.

2.3 Use cases by robot family

The same platform serves several kinds of maker, and the design study draws scenarios from each.

Robot family	Representative goal	World it is validated in
Rover	Reach a marker over rough open ground	Open terrain, then planetary worlds
Self driving car	Cross a junction without hitting a pedestrian or a vehicle	Riverside City, among moving traffic
Personal companion	Move through a furnished room and fetch an object	The indoor Living Room
Robotic arm	Pick an object up and set it down	The indoor Living Room, at a table
Custom microcontroller build	Whatever the maker sets, on parts they chose	Recommended by the assistant from the build

3. Design

The platform is written in Python 3.12 and presented through a desktop web view, with the interface built in React and pre-compiled to a single offline bundle, the three dimensional world drawn with a vendored copy of Three.js, and the user's designs and the system's accumulated experience stored in a single local SQLite file. The artificial intelligence calls a small open weight model served locally by Ollama. The robot the user designs is the single source of truth the rest of the system reads: its board, sensors and motors set the mass, the top speed and the commands that exist, and the grounded assistant reads the same specification so it cannot suggest a part the build does not have. Behaviour is validated across randomised seeds after Tobin et al. (2017), in worlds chosen to suit the robot, from an open terrain to a street with pedestrians and traffic to an indoor room.

3.1 Self refinement design

Self refinement is delivered at the level of the system and is described as refinement, not improvement, on purpose. Three mechanisms carry it, all local and all measurable. After each run the system writes a short reflection on what failed and what worked and retrieves the most similar past reflections into the next suggestion, after Shinn et al. (2023). Control routines that passed validation are kept as named, parameterised skills in a growing library and composed into later designs, after Wang et al. (2023). An experience table of past designs and their outcomes is sampled so advice on a new build is anchored in what actually happened on similar ones. None of this edits the model.

3.2 Architecture and the central loop

The system is organised around a single loop the interface makes concrete: design, program, validate, refine. The Robot Lab produces a specification; the editor produces a program; the simulator runs that program against a scenario in a chosen world and produces a result; and the memory turns that result into a reflection and, when the program worked, a reusable skill. The specification is the hub the other parts read, which is what keeps the assistant honest, since it can only speak about the build that exists. A thin Python layer hosts the simulation engine, the sandbox and the bridge to the local model, while the interface, the worlds and the agents run in the desktop web view; the two communicate over a small, typed bridge rather than sharing state, so each can be tested on its own.




Figure 3.1. The central loop. The user designs a robot, programs it, validates it in simulation, and the system refines from the result; the build specification is the hub every part reads, and the experience memory carries lessons from Refine back into the next Design.

3.3 The world and the agents

A world is described by its environment, its terrain, a field of fixed obstacles and a set of moving agents. The environment sets gravity, temperature, light and traction, which feed the telemetry and the physics; the fixed obstacles, such as buildings and parked cars in the city or furniture in the room, are what the robot must not hit; and the agents, pedestrians and traffic, move with a will of their own. The agent design follows the social force model of pedestrian motion (Helbing and Molnar, 1995) and reciprocal velocity obstacles for mutual avoidance (van den Berg et al., 2008); the realised system uses a simplified, deterministic form of these so the same scene replays identically, with the full force based model named as the next step. A single agent simulation is the one source of truth for every consumer: the three dimensional view, the flat view and the robot's collision and distance sensor all read the same positions, so an agent the robot can see is one it can hit.

3.4 Choosing the world for a robot

A robot is only meaningfully tested where it would actually work, so the assistant reasons from the build to a world: a road vehicle is sent to the city among traffic, a companion or an arm to the indoor room, an explorer to open terrain. This is deliberately practical rather than decorative; placing a car in a living room teaches nothing, so the recommendation is shown with its reason and the new robot is dropped straight into the world that suits it, with the others available for the user to try next.

4. Implementation

This chapter describes the realised system. The headline components built to date are the Robot Lab, where a user designs a custom robot whose specification drives the simulation; the realistic worlds, including Riverside City with road markings, a crossing, buildings, pedestrians and traffic; the grounded assistant, the second agent code reviewer and the voice route; the multi robot swarm; and the offline self refinement store.

4.1 The Robot Lab and the specification

The Robot Lab is where a build becomes data. The user picks an archetype, a controller board such as an ESP32 or a micro:bit, a set of sensors and a set of actuators, and the panel computes a specification as it goes: the total mass from the chassis and the chosen parts, an estimated battery runtime, a top speed from the strongest motor set, and the list of commands the build supports, where each sensor unlocks a matching call and an absent sensor withholds it. This specification is not cosmetic. The simulation reads it on every move, so a heavier build drains its battery faster and a stronger motor set drives quicker; and the assistant reads the same specification, so it cannot suggest a command for a part the build does not carry. The robot rendered in the world is built to match the archetype as well, so a car looks like a car and a companion like a companion, which makes the design tangible rather than abstract.




Figure 4.1. The Robot Lab. Choosing a board, sensors and actuators produces a live specification, here a self-driving car, with mass, battery life, top speed and the supported commands updating as the build changes.

4.2 The three dimensional worlds

The worlds are drawn with a vendored copy of the Three.js library held in the project, so the graphics work offline. Each world is a real scene rather than a backdrop. The city is built from crossing roads with lane markings and a pedestrian crossing, buildings with lit window textures, and parked and moving vehicles; the indoor room is a furnished space with walls, a floor, a sofa,

a table on a rug and a shelf, lit warmly for an interior. The renderer uses soft shadows and filmic tone mapping so the result reads as a place with depth that the user can orbit through three hundred and sixty degrees, rather than a flat diagram. The same worlds are also offered in a lighter flat view for machines without a usable graphics card, drawn from the same underlying world description so the two views show the same place; the robot is scaled to fit indoors so it shares the room with the furniture rather than towering over it.




Figure 4.2. Riverside City, rendered in real three dimensional graphics: crossing roads with markings and a pedestrian crossing, buildings with lit windows, and vehicles and pedestrians that move around the robot.



Figure 4.3. The indoor Living Room for companion and humanoid robots, with the robot sized to share the space with the furniture and the people who move through it.

4.3 The moving agent simulation

Pedestrians and traffic are driven by one small simulation that is the single source of truth for the whole system. It holds each agent's position in the same world coordinates the robot and the obstacles use, advances them along their paths each frame, and exposes the current positions to every consumer. The three dimensional view renders a mesh per agent at those positions, the flat view renders a sprite, and crucially the robot's collision test and its distance sensor read the very same list, so a pedestrian the robot can see is a pedestrian it can hit and a sensor can detect. This unification is what makes the realistic world more than scenery: the robot genuinely has to cope with what moves around it.

4.4 The artificial intelligence layer

The assistant runs entirely on a small local model served by Ollama, with no network call. Every prompt is grounded in the robot's specification so the model is constrained to the build's real capabilities, which is the main defence against the model inventing a command that does not exist. A second agent reviews generated control code before it is offered, and nothing a user could run reaches a real robot without first running in the sandboxed simulation. When no model is installed the same functions return deterministic, rule based code and checks within about two seconds, so the platform stays fully usable offline and on weak hardware; the model adds conversational phrasing and flexibility on top of a path that always works.

4.5 The self refinement store

Self refinement is implemented as a local store, not as any change to the model, and it is the honest, countable form of the idea. After each run the system records the outcome and writes a

short reflection on what happened, in the spirit of verbal reflection (Shinn et al., 2023); a run that crashes into a pedestrian yields a reflection about slowing and sensing before crossing, and a clean finish yields one about saving the working program. Programs that worked can be kept as named skills in a growing library and reused on later designs, after the skill library idea of Wang et al. (2023). Both the reflections and the skills are surfaced in a memory panel and persist in the local file across sessions, so the system carries forward what it has seen without ever retraining the model (Tao et al., 2024). The verified behaviour of this loop, a designed car driving the city, the distance sensor picking up the agents, a collision recorded and the reflection appearing in the panel, is reported in the evaluation.

4.6 The sandbox and the verifier

A pupil's, or a maker's, program is real Python, so it runs in a sandbox that blocks file, network and process access, and its behaviour is graded against the scenario's success criteria rather than trusted. A design is run many times with friction, mass and sensor noise randomised after Tobin et al. (2017) and judged on the spread, not a single tidy run, which is what lets the platform speak honestly about how a robot would behave rather than how it behaved once.

4.7 Engineering, build and packaging

The work is organised as small, independently testable changes, each covered by automated tests and merged only once a continuous integration pipeline runs the full test suite, a type checker and a linter on three operating systems, so the codebase stays releasable at every step. The interface is pre-compiled from its React and JSX sources into a single offline bundle by a vendored compiler, so the shipped application loads plain JavaScript rather than transpiling on every cold start. The result is delivered as a single windowed desktop application, built with PyInstaller, that installs without administrator rights and runs the modern interface in a native web view with the local model fetched once through Ollama; the build was produced and run as part of this work, confirming the single download claim is real rather than aspirational.

5. Testing and Evaluation

Evaluation asks three separate questions. The first is whether the platform is usable, tested continuously with a simulated persona panel in the tradition of discount usability inspection (Nielsen, 1994). The second is whether it helps: whether a non-expert using Kodro produces a robot that behaves better in a held out realistic scenario than one designed without it. The third is whether the artificial intelligence genuinely self refines, measured directly by whether the median iterations to a working design fall as the experience memory grows, the skill library's reuse rate rises, and a retrieved reflection more often precedes a successful fix than chance would give. The artificial intelligence is also measured on its own terms, with success bars set in advance: grounded answers correct at least eighty five percent of the time and inventing content less than five percent of the time, the latter a release gate.

5.1 Automated testing

Testing is automated and continuous. The project carries more than eight hundred unit and integration tests at around eighty five percent line coverage, with a type checker and a linter enforced on every change across three operating systems, so a regression is caught as a failing build rather than by a user. This is the floor the rest of the evaluation stands on: the worlds, the sandbox and the bridge are exercised by tests before any human judgement is asked for.

5.2 Usability: the simulated persona panel

Usability was tested continuously with a simulated persona panel, a low cost inspection method in the tradition of discount usability (Nielsen, 1994), where a set of distinct personas spanning the makers who might use the tool each drives the current build and rates it out of ten, naming the concrete defects it found. After each round the named defects were fixed and the same build was re rated, so the trend reflects real change in the product. The trajectory is set out below; it is reported as a usability finding only, and says nothing about whether anyone learns or builds a better robot, which the design study alone answers.

Round	Personas	Focus	Mean rating (out of 10)
1	50	First contact with the build	5.86
2	12	After the first round of fixes	6.50
3	50	After the learner model and maker tools landed	7.10
4	50	After grounded answers and the voice route	7.36
5	8	Focused review of the new three dimensional world	6.40

The mean rose from about 5.9 at first contact to about 7.4 on the core product as named defects were fixed between rounds. The fifth round was a fresh, focused inspection of the newly added three dimensional view, which is why it sits lower; the issues it raised, around readable direction, reduced motion and performance on weak graphics hardware, were then addressed in the same

loop. The generation prompt and rating scale are reproduced in Appendix B so the panel can be regenerated.

5.3 The self refinement loop, observed

The self refinement mechanism is implemented, and its end to end behaviour was observed directly during development: a car designed in the Robot Lab was driven into the city, the distance sensor registered the moving agents ahead, a collision was recorded, and a written reflection on that outcome appeared in the memory panel and persisted across the session. This confirms the loop works as designed; it is not yet the controlled ablation, which is set out next.

5.4 Studies still to run

To fill with measured results: the design study (within subject, about sixteen makers, with and without the tool, task success over twenty randomised trials, iterations to a working design, and the before and after behaviour prediction and code quality learning measures); the self refinement ablation (memory on versus off, against a rule only baseline, with the stated effect size); and the artificial intelligence honesty figures (correctness, invention and refusal rates, including under adversarial prompts). All figures will be the measured numbers, weak results included.

6. Project Ethics

The declared categories are data category Ao and participant category 2. The development and the simulated persona evaluation involved no human participants and no personal data. The design study with human participants used adult and older makers and students rather than young children, which kept the safeguarding burden proportionate, with plain explanation, consent, and anonymisation at the point of collection. One risk is specific to the technology and deserves naming on its own: the output here is control code for a machine a user may later build, so a confident falsehood could produce code that drives a real robot into a person or a wall. The design treats this as a safety requirement: every suggestion is grounded in the robot's specification, the assistant refuses rather than guesses, and any code runs first in the sandboxed simulation against the scenario, so a wrong suggestion fails visibly on screen rather than on a real robot. The mandatory simulation gate, not the model, stands between a generated program and the real world.

7. Conclusion and Future Work

This project set out to build an offline platform on which a non expert could design a robot, program it, and validate how it behaves in a realistic simulation, and to ask how far an artificial intelligence assistant could usefully refine itself under a hard constraint that ruled out the cloud and ruled out changing the model's weights. The platform was built and runs as a single installed application: a Robot Lab whose specification drives the simulation, three dimensional worlds from a street with traffic to an indoor room, a moving agent simulation the robot genuinely has to cope with, an assistant grounded in the build with a deterministic fallback, and a self refinement store of reflections and reusable skills whose loop was observed working end to end.

On the central question, the honest answer is bounded and was bounded on purpose. Weight level self evolution remains unsolved and was never attempted; what the system does instead is refine at the level of the system, remembering and reusing what worked, which is countable and was shown to function. The validation is similarly honest: a simulator with randomised conditions can rehearse how a robot would behave and surface failures cheaply, but it cannot promise the real world, and the work says so plainly rather than overclaiming. The strongest evidence in hand is the usability trajectory and the working software; the questions of whether a maker learns and whether the refinement measurably helps are designed in detail but await the field study and the ablation.

Future work falls into three lines. The first is fidelity: richer three dimensional assets through an offline modelling pipeline, and dynamic agents that avoid the robot and each other with the full social force and reciprocal velocity models the design already names. The second is breadth: more robot families and more indoor and outdoor worlds, and wiring the stored reflections back into the assistant's prompt so past lessons shape new advice directly. The third is evidence: running the design study and the self refinement ablation at a larger scale than a single project allows, to turn the designed measures into measured results.

References

- Ahn, M., Brohan, A., Brown, N. et al. (2022) Do as I can, not as I say: grounding language in robotic affordances. *arXiv:2204.01691*.
- Helbing, D. and Molnar, P. (1995) Social force model for pedestrian dynamics. *Physical Review E*, 51(5), pp. 4282 to 4286.
- Huang, L., Yu, W., Ma, W. et al. (2023) A survey on hallucination in large language models. *ACM Transactions on Information Systems* (arXiv:2311.05232).
- Lewis, P., Perez, E., Piktus, A. et al. (2020) Retrieval augmented generation for knowledge intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33.
- Liang, J., Huang, W., Xia, F. et al. (2023) Code as policies: language model programs for embodied control. *IEEE International Conference on Robotics and Automation (ICRA)*.
- Nielsen, J. (1994) Usability engineering. San Francisco: Morgan Kaufmann.
- Papert, S. (1980) Mindstorms: children, computers, and powerful ideas. New York: Basic Books.
- Shinn, N., Cassano, F., Berman, E. et al. (2023) Reflexion: language agents with verbal reinforcement learning. *arXiv:2303.11366*.
- Sweller, J. (1988) Cognitive load during problem solving: effects on learning. *Cognitive Science*, 12(2), pp. 257 to 285.
- Tao, Z., Lin, T.E., Chen, X. et al. (2024) A survey on self evolution of large language models. *arXiv:2404.14387*.
- Tobin, J., Fong, R., Ray, A. et al. (2017) Domain randomization for transferring deep neural networks from simulation to the real world. *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pp. 23 to 30.
- van den Berg, J., Lin, M. and Manocha, D. (2008) Reciprocal velocity obstacles for real-time multi-agent navigation. *IEEE International Conference on Robotics and Automation (ICRA)*, pp. 1928 to 1935.
- Wang, G., Xie, Y., Jiang, Y. et al. (2023) Voyager: an open-ended embodied agent with large language models. *arXiv:2305.16291*.

Appendices

Appendix A. Code archive. The full source, build scripts and test suite accompany this dissertation as supplementary material.

Appendix B. Persona generation. The persona generation prompt and rating scale used in the usability evaluation are reproduced here.

To attach: the code archive, the persona prompt, additional evaluation tables, and any material that supports but is not essential to the main 50 pages.

Working draft. Generative artificial intelligence assisted the development of the software and the drafting of this document, acknowledged in line with University guidance. Sections marked to complete are filled as the build and evaluation finish; all reported figures will be the measured results.